

Assignment No 2

Problem Statement -

Perform the following operations using Python on the Air quality data sets

- a. Data cleaning
- b. Data integration
- c. Data transformation
- d. Error correcting
- e. Data model building

Importing required libraries

```
In [58]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import warnings
warnings.filterwarnings('ignore')
import seaborn as sns
```

Reading dataset

```
In [59]: data = pd.read_csv("airquality2.csv")
data
```

Out[59]:

	Unnamed: 0	Ozone	Solar.R	Wind	Temp	Month	Day	Humidity
0	1	41.0	190.0	7.4	67	5	1	High
1	2	36.0	118.0	8.0	72	5	2	High
2	3	12.0	149.0	12.6	74	5	3	NaN
3	4	18.0	313.0	11.5	62	5	4	Medium
4	5	NaN	NaN	14.3	56	5	5	Low
...
148	149	30.0	193.0	6.9	70	9	26	Low
149	150	NaN	145.0	13.2	77	9	27	Low
150	151	14.0	191.0	14.3	75	9	28	High
151	152	18.0	131.0	8.0	76	9	29	High
152	153	20.0	223.0	11.5	68	9	30	Medium

153 rows × 8 columns

```
In [60]: data.drop(data.iloc[:,[0]], axis=1, inplace=True)
data
```

Out[60]:

	Ozone	Solar.R	Wind	Temp	Month	Day	Humidity
0	41.0	190.0	7.4	67	5	1	High
1	36.0	118.0	8.0	72	5	2	High
2	12.0	149.0	12.6	74	5	3	NaN
3	18.0	313.0	11.5	62	5	4	Medium
4	NaN	NaN	14.3	56	5	5	Low
...
148	30.0	193.0	6.9	70	9	26	Low
149	NaN	145.0	13.2	77	9	27	Low
150	14.0	191.0	14.3	75	9	28	High
151	18.0	131.0	8.0	76	9	29	High
152	20.0	223.0	11.5	68	9	30	Medium

153 rows × 7 columns

```
In [61]: data.isnull().sum()
```

```
Out[61]: Ozone      37
          Solar.R    7
          Wind       0
          Temp       0
          Month      0
          Day        0
          Humidity    9
          dtype: int64
```

```
In [62]: data["Ozone"].fillna(data["Ozone"].mean(), inplace=True)
          data["Solar.R"].fillna(data["Solar.R"].mean(), inplace=True)
          data
```

```
Out[62]:
```

	Ozone	Solar.R	Wind	Temp	Month	Day	Humidity
0	41.00000	190.000000	7.4	67	5	1	High
1	36.00000	118.000000	8.0	72	5	2	High
2	12.00000	149.000000	12.6	74	5	3	NaN
3	18.00000	313.000000	11.5	62	5	4	Medium
4	42.12931	185.931507	14.3	56	5	5	Low
...
148	30.00000	193.000000	6.9	70	9	26	Low
149	42.12931	145.000000	13.2	77	9	27	Low
150	14.00000	191.000000	14.3	75	9	28	High
151	18.00000	131.000000	8.0	76	9	29	High
152	20.00000	223.000000	11.5	68	9	30	Medium

153 rows × 7 columns

```
In [63]: data["Humidity"].fillna("Medium", inplace=True)
```

```
In [64]: data.isnull().sum()
```

```
Out[64]: Ozone      0
          Solar.R    0
          Wind       0
          Temp       0
          Month      0
          Day        0
          Humidity    0
          dtype: int64
```

Performing Label Encoding

```
In [65]: from sklearn.preprocessing import LabelEncoder
```

```
In [66]: le = LabelEncoder()
```

```
In [67]: data["Humidity"] = le.fit_transform(data["Humidity"])
```

```
In [68]: data
```

Out[68]:

	Ozone	Solar.R	Wind	Temp	Month	Day	Humidity
0	41.00000	190.000000	7.4	67	5	1	0
1	36.00000	118.000000	8.0	72	5	2	0
2	12.00000	149.000000	12.6	74	5	3	2
3	18.00000	313.000000	11.5	62	5	4	2
4	42.12931	185.931507	14.3	56	5	5	1
...
148	30.00000	193.000000	6.9	70	9	26	1
149	42.12931	145.000000	13.2	77	9	27	1
150	14.00000	191.000000	14.3	75	9	28	0
151	18.00000	131.000000	8.0	76	9	29	0
152	20.00000	223.000000	11.5	68	9	30	2

153 rows × 7 columns

Assigning dependent and independent variables

```
In [69]: x = data.iloc[:, [0,1,2,3]]
         y = data["Humidity"]
```

Importing the required function and splitting the data into training and testing data

```
In [70]: from sklearn.model_selection import train_test_split
```

```
In [71]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2)
```

Printing size of training and testing data

```
In [72]: print(len(x_train))
print(len(x_test))
print(len(y_train))
print(len(y_test))
```

```
122
31
122
31
```

Importing and creating a linear regression model

```
In [73]: from sklearn import linear_model
```

```
In [74]: regr = linear_model.LinearRegression()
```

Fitting the model and display the regression coefficients

```
In [75]: model = regr.fit(x_train, y_train)
print(model.intercept_)
print(model.coef_)
```

```
0.2596553365337608
[-0.00196037 -0.00092793  0.02911026  0.01001252]
```

Predicting the values

```
In [76]: y_predict = model.predict(x_test)
```

```
In [77]: y_predict
```

```
Out[77]: array([1.07433468, 1.09123502, 0.79142388, 1.13114234, 0.94300376,
                1.11784014, 1.09664385, 0.80813102, 0.89763294, 1.12174905,
                0.9709549 , 1.00403062, 1.07364943, 1.2003274 , 1.10474293,
                1.16587009, 1.07330507, 1.17589964, 1.24793266, 0.99615301,
                0.9414399 , 1.18835405, 1.04807664, 1.19328169, 0.97611117,
                0.70579647, 1.01100268, 0.90005967, 1.19916514, 1.03337014,
                1.06991779])
```

Importing and calculating the performance metrics

```
In [78]: from sklearn.metrics import mean_squared_error
```

Mean Squared Error (MSE)

```
In [79]: mse = mean_squared_error(y_test,y_predict)
mse
```

```
Out[79]: 0.7555852807566938
```


Root Mean Squared Error (RMSE)

```
In [80]: rmse = np.sqrt(mse)
rmse
```

```
Out[80]: 0.8692440858336017
```

Plotting the data

```
In [83]: plt.scatter(y_test,y_predict);  
plt.xlabel('Actual');  
plt.ylabel('Predicted');
```



```
In [85]: sns.regplot(x=y_test,y=y_predict,ci=None,color='red');
```

